

IRIS: Innovation Residual Iterative Stabilization

Learning Rate Robustness through Error-Correcting Optimization

Udit Asthana
TravVir
madhavasthana@gmail.com

July 25, 2026

Abstract

Adaptive gradient methods face a fundamental tradeoff: aggressive learning rates accelerate convergence but risk instability, while conservative rates require extensive tuning. We introduce **IRIS** (Innovation Residual Iterative Stabilization), an optimizer that tolerates learning rates **10-1000× higher** than existing methods through innovation-based error correction. Unlike gradient difference methods (e.g., Adan: $d_t = \mathbf{g}_t - \mathbf{g}_{t-1}$) which exhibit $2\times$ variance and convergence-divergence cycles, IRIS tracks *innovation*—prediction error relative to smooth momentum ($\mathbf{I}_{t|t-1} = \mathbf{g}_t - \hat{\mathbf{g}}_{t-1}$)—enabling automatic stabilization even with extreme learning rates.

On toy optimization landscapes, IRIS achieves 10-1000× better convergence than Adam-family optimizers (Rosenbrock with LR=1.0: IRIS reaches distance 0.07, Adam fails at LR>0.01) and uniquely escapes local minima on Rastrigin where all competitors fail. On CIFAR-100 (ResNet-18, batch 2048), IRIS achieves +0.8% accuracy over AdamW while training 25% faster (113 vs 150 minutes) with no warmup required. IRIS maintains Adam-equivalent computational cost while eliminating hyperparameter sensitivity: a single learning rate (0.3-1.0) works universally across tasks.

Code: <https://github.com/madhavasthana-oss/iris>

1 Introduction

Deep learning’s success hinges on efficient optimization, yet adaptive gradient methods remain surprisingly brittle. Adam [Kingma and Ba, 2015] and its variants require extensive hyperparameter tuning—particularly learning rate—and exhibit training instabilities that necessitate multi-epoch warmup and carefully designed schedules. Recent attempts to address these limitations introduce new pathologies: AdaBelief [Zhuang et al., 2020] catastrophically overfits in large-batch settings, while Adan [Xie et al., 2024] exhibits convergence-divergence cycles that make training unpredictable.

1.1 The Core Problem: Instability vs Convergence Speed

Current optimizers face an unavoidable tradeoff:

- **Adam/AdamW:** Stable but requires $\text{LR} \leq 3e-3$, extensive warmup (20 epochs), and cannot distinguish stochastic noise from genuine landscape difficulty
- **AdaBelief:** Uses prediction error $(\mathbf{g}_t - \mathbf{m}_t)^2$ in denominator, but this “belief” is corrupted (β_1 -scaled innovation), causing severe overfitting at large batch sizes
- **Adan:** Uses gradient difference $d_t = \mathbf{g}_t - \mathbf{g}_{t-1}$ which has $2\times$ variance (comparing two noisy samples), leading to perpetual oscillations and mid-training catastrophic spikes

1.2 Our Approach: Innovation-Based Error Correction

We introduce IRIS, which frames optimization as **predictive error correction** rather than gradient descent or momentum extrapolation. IRIS tracks *innovation*—the prediction error $\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1}$ —and uses it to:

1. **Correct gradient estimates:** $\widehat{\mathbf{g}}_t = \widehat{\mathbf{g}}_{t-1} + \psi_1^{-1} \mathbf{I}_{t|t-1}$ (Kalman-style update)
2. **Accumulate systematic errors:** $\widehat{\boldsymbol{\varepsilon}}_t = \widehat{\boldsymbol{\varepsilon}}_{t-1} + \psi_2^{-1} (\mathbf{I}_{t|t-1} - \widehat{\boldsymbol{\varepsilon}}_{t-1})$
3. **Match variance to updates:** $\widehat{\mathbf{v}}_t \propto (\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2$ (variance of corrected gradient)

Key innovation: Unlike Adan’s gradient difference (compares two noisy samples: $\text{Var}[d_t] = 2\sigma^2$), IRIS’s innovation compares current gradient to smooth estimate ($\text{Var}[\mathbf{I}_{t|t-1}] = \sigma^2$), reducing variance by 50% and enabling hyperstable convergence.

1.3 Main Contributions

1. **Theoretical:** We prove innovation-based methods have 50% lower variance than gradient difference methods (Proposition 1), enabling 10-1000× higher learning rate tolerance
2. **Algorithmic:** We introduce recursive bias correction via Kalman gains ($\psi_{i,t} = \beta_i \psi_{i,t-1} + 1$), eliminating GPU-CPU synchronization and enabling exact bias correction from step 1 (no warmup needed)
3. **Empirical:** We demonstrate:
 - **Extreme LR robustness:** Rosenbrock with LR=1.0 (IRIS: 0.07 distance, Adan fails at LR>0.15, Adam at LR>0.01)
 - **Local minima escape:** Rastrigin global minimum found (distance 10^{-6}), all competitors trapped at ~ 1
 - **Real-world gains:** CIFAR-100 batch 2048: +0.8% accuracy, 25% faster (113 vs 150 min), no warmup
4. **Practical:** IRIS requires *no hyperparameter tuning*—a single LR (0.3-1.0) works universally, eliminating the grid search that plagues existing optimizers

1.4 Paper Organization

Section 2 reviews adaptive optimization and identifies limitations. Section 3 introduces IRIS’s innovation-based framework and contrasts it with gradient difference methods. Section 4 provides theoretical analysis including variance reduction proofs and convergence guarantees. Section 5 presents comprehensive benchmarks on toy functions and real tasks. Section 6 discusses implications and limitations.

2 Related Work

2.1 Adaptive Gradient Methods

First-order adaptive methods like Adam [Kingma and Ba, 2015], RMSprop [Tieleman and Hinton, 2012], and AdaGrad [Duchi et al., 2011] adapt learning rates per parameter using gradient statistics. AdamW [Loshchilov and Hutter, 2019] improved Adam by decoupling weight decay, while AMSGrad [Reddi et al., 2018] addressed convergence issues via maximum variance tracking. However, all share fundamental limitations: (1) conflating noise with curvature in second moments, (2) narrow learning rate ranges, (3) warmup requirements.

AdaBelief [Zhuang et al., 2020] attempted to measure “prediction error” via $(\mathbf{g}_t - \mathbf{m}_t)^2$, but this quantity is corrupted: it equals $\beta_1(\mathbf{g}_t - \mathbf{m}_{t-1})$ at steady state—a scaled, biased version of innovation. Critically, AdaBelief uses this term *inversely* (in denominator), causing division-by-zero instabilities and severe overfitting in large-batch settings (our CIFAR-100 experiments: -2% accuracy vs IRIS despite lower training loss).

2.2 Momentum and Extrapolation Methods

Nesterov momentum [Nesterov, 1983] uses gradient at extrapolated position. **Adan** [Xie et al., 2024] extends this with gradient difference $d_t = \mathbf{g}_t - \mathbf{g}_{t-1}$ for adaptive Nesterov momentum. However, gradient difference has fundamental flaws:

- **High variance:** $\text{Var}[d_t] = \text{Var}[\mathbf{g}_t] + \text{Var}[\mathbf{g}_{t-1}] \approx 2\sigma^2$ (two noisy samples)
- **Vanishing correction:** Near minima, $\mathbf{g}_t \approx \mathbf{g}_{t-1} \implies d_t \approx 0$ (no stabilization)
- **Convergence-divergence cycles:** Our experiments show catastrophic mid-training spikes (Rosenbrock: $10^{-4} \rightarrow 10^{-1}$ at step 400; Rastrigin: repeated $10^6\times$ spikes)

IRIS addresses these issues by using innovation $\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1}$ (gradient vs smooth estimate), reducing variance to σ^2 and maintaining active correction even when gradients vanish.

2.3 Second-Order Methods

Newton-type methods like L-BFGS [Liu and Nocedal, 1989], K-FAC [Martens and Grosse, 2015], and Shampoo [Gupta et al., 2018] use curvature information but incur prohibitive computational costs ($O(d^2)$ or worse). **AdaHessian** [Yao et al., 2021] approximates diagonal Hessian efficiently but still requires Hessian-vector products. **Sophia** [Liu et al., 2024] clips by Hessian diagonal but exhibits aggressive clipping (70% engagement) and $2\times$ per-step cost.

IRIS achieves second-order-like adaptivity at first-order cost: innovation variance $(\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2$ implicitly captures curvature without explicit Hessian computation.

2.4 Kalman Filtering in Optimization

Extended Kalman filters [Kalman, 1960] provide optimal state estimation under Gaussian noise. Recent work on Kalman-inspired optimizers [Vuckovic, 2018, Olivier et al., 2021] focused primarily on state estimation rather than error correction. IRIS’s key innovation is using the *innovation sequence* (prediction errors) as a first-class signal for gradient correction and variance estimation, inspired by Kalman filtering’s innovation analysis but adapted to stochastic optimization’s unique challenges.

3 Innovation-Based Error Correction

3.1 Motivation: Why Innovation Over Gradient Difference?

Consider optimizing $f(\mathbf{x})$ with stochastic gradients $\mathbf{g}_t = \nabla f(\mathbf{x}_t) + \xi_t$ where $\mathbb{E}[\xi_t] = 0$, $\text{Var}[\xi_t] = \sigma^2$.

Adan’s gradient difference:

$$d_t = \mathbf{g}_t - \mathbf{g}_{t-1} \tag{1}$$

This compares *two noisy samples*. Under weak correlation:

$$\text{Var}[d_t] = \text{Var}[\mathbf{g}_t] + \text{Var}[\mathbf{g}_{t-1}] - 2\text{Cov}[\mathbf{g}_t, \mathbf{g}_{t-1}] \approx 2\sigma^2 \tag{2}$$

IRIS's innovation:

$$\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1} \quad (3)$$

This compares *noisy sample* to *smooth estimate* $\widehat{\mathbf{g}}_{t-1} = \beta_1 \widehat{\mathbf{g}}_{t-2} + (1 - \beta_1) \mathbf{g}_{t-1}$:

$$\text{Var}[\mathbf{I}_{t|t-1}] = \text{Var}[\mathbf{g}_t] + \underbrace{\text{Var}[\widehat{\mathbf{g}}_{t-1}]}_{\approx 0} \approx \sigma^2 \quad (4)$$

Proposition 1 (Variance Reduction). *For stochastic gradients with variance σ^2 and weak temporal correlation, gradient difference has approximately 2× the variance of innovation:*

$$\frac{\text{Var}[d_t]}{\text{Var}[\mathbf{I}_{t|t-1}]} \approx 2 \quad (5)$$

Implication: Innovation's 50% lower variance enables stable convergence even with high variance smoothing ($\beta_3 \geq 0.999$), allowing 10-1000× higher learning rates.

3.2 The IRIS Algorithm

IRIS maintains three bias-corrected estimates:

- $\widehat{\mathbf{g}}_t$: Gradient estimate (momentum)
- $\widehat{\mathbf{e}}_t$: Innovation residual (systematic error accumulation)
- $\widehat{\mathbf{v}}_t$: Variance estimate (uncertainty of corrected gradient)

3.3 Key Design Choices

3.3.1 Variance Matching (Line 9)

The variance term $(\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2$ estimates uncertainty of the *corrected gradient*—exactly what appears in the numerator. This ensures the preconditioner scale matches the update scale.

Contrast with AdaBelief: AdaBelief uses $(\mathbf{g}_t - \mathbf{m}_t)^2$ in the *denominator*, causing:

- Division by small numbers when $\mathbf{g}_t \approx \mathbf{m}_t$ (instability)
- Inverse scaling: small prediction error $\implies$ large step (counterintuitive)
- No variance matching: denominator doesn't match numerator structure

IRIS uses innovation in the *numerator*: correction modifies *direction*, not scale. Variance matching in denominator provides stable preconditioning.

3.3.2 Recursive Bias Correction (Line 4)

Standard Adam uses $\hat{m}_t = m_t / (1 - \beta_1^t)$, requiring:

- Integer t from GPU (GPU-CPU synchronization)
- Exponentiation β_1^t (numerical instability at large t)

Algorithm 1 IRIS: Innovation Residual Iterative Stabilization

Require: Learning rate η , weight decay λ , numerical stability ε

Require: EMA coefficients $(\beta_1, \beta_2, \beta_3)$, optional split-correction γ

Initialize $\hat{\mathbf{g}}_0 \leftarrow \mathbf{0}$, $\hat{\varepsilon}_0 \leftarrow \mathbf{0}$, $\hat{\mathbf{v}}_0 \leftarrow \mathbf{0}$

Initialize bias correction $\psi_{1,0} \leftarrow 0$, $\psi_{2,0} \leftarrow 0$, $\psi_{3,0} \leftarrow 0$

for $t = 1, 2, 3, \dots$ **do**

$\psi_{i,t} \leftarrow \beta_i \cdot \psi_{i,t-1} + 1$, $\forall i \in \{1, 2, 3\}$

// Bias correction accumulators

if split-correction γ **then**

$\varphi_t \leftarrow \gamma \cdot \varphi_{t-1} + 1$, $\forall i \in \{1, 2, 3\}$

// Bias correction accumulators

$\mathbf{I}_{t|t-1} \leftarrow \mathbf{g}_t - \hat{\mathbf{g}}_{t-1}$

// Innovation (prediction error)

$\hat{\mathbf{g}}_t \leftarrow \hat{\mathbf{g}}_{t-1} + \psi_{1,t}^{-1} \mathbf{I}_{t|t-1}$

// Update gradient estimate

$\hat{\varepsilon}_t \leftarrow (1 - \psi_{2,t}^{-1}) \hat{\varepsilon}_{t-1} + \psi_{2,t}^{-1} \mathbf{I}_{t|t-1}$

$\hat{\mathbf{v}}_t \leftarrow (1 - \psi_{3,t}^{-1}) \hat{\mathbf{v}}_{t-1} + \psi_{3,t}^{-1} \mathbf{g}_t^2$

// Update gradient estimate

$\hat{\mathbf{n}}_t \leftarrow (1 - \varphi_t^{-1}) \hat{\mathbf{n}}_{t-1} + \varphi_t^{-1} \mathbf{I}_{t|t-1}^2$

$\mathbf{x}_t \leftarrow (1 - \eta\lambda) \mathbf{x}_{t-1} - \eta \left(\frac{\hat{\mathbf{g}}_t}{\sqrt{\hat{\mathbf{v}}_t} + \varepsilon} + \beta_2 \frac{\hat{\varepsilon}_t}{\sqrt{\hat{\mathbf{n}}_t} + \varepsilon} \right)$

// Parameter update

else

$\mathbf{I}_{t|t-1} \leftarrow \mathbf{g}_t - \hat{\mathbf{g}}_{t-1}$

// Innovation (prediction error)

$\mathbf{E}_{t|t-1} \leftarrow \mathbf{I}_{t|t-1} - \hat{\varepsilon}_{t-1}$

// Error-of-innovation

$\Sigma_{t|t-1} \leftarrow (\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2 - \hat{\mathbf{v}}_{t-1}$

// Variance innovation

$\hat{\mathbf{g}}_t \leftarrow \hat{\mathbf{g}}_{t-1} + \psi_{1,t}^{-1} \mathbf{I}_{t|t-1}$

// Update gradient estimate

$\hat{\varepsilon}_t \leftarrow \hat{\varepsilon}_{t-1} + \psi_{2,t}^{-1} \mathbf{E}_{t|t-1}$

// Update innovation residual

$\hat{\mathbf{v}}_t \leftarrow \hat{\mathbf{v}}_{t-1} + \psi_{3,t}^{-1} \Sigma_{t|t-1}$

// Update variance

$\mathbf{x}_t \leftarrow (1 - \eta\lambda) \mathbf{x}_{t-1} - \eta \frac{\hat{\mathbf{g}}_t + \beta_2 \hat{\varepsilon}_t}{\sqrt{\hat{\mathbf{v}}_t} + \varepsilon}$

// Parameter update

end if

end for

return $\mathbf{x}_t$

IRIS uses recursive accumulation:

$$\psi_{i,t} = \beta_i \psi_{i,t-1} + 1 \implies \psi_{i,t}^{-1} = \frac{\beta_i^{-1}}{\psi_{i,t-1}^{-1} + \beta_i^{-1}} \quad (6)$$

This is equivalent to Kalman gain with static measurement noise $R = \beta_i^{-1}$:

$$K_t = \frac{P_{t-1}}{P_{t-1} + R} \quad (\text{Kalman filter}) \quad (7)$$

where $P_t \leftrightarrow \psi_t^{-1}$ (accumulated uncertainty), $R \leftrightarrow \beta^{-1}$ (measurement noise).

Benefits: No GPU-CPU sync, no exponentiation, exact bias correction from step 1 $\implies$ **no warmup required.**

3.3.3 Timescale Separation ($\beta_2 < \beta_1$)

We empirically find optimal $\beta_2 = \beta_1 - \delta$ where $\delta \in [0.15, 0.25]$. This creates timescale separation:

$$\tau_{\text{gradient}} = \frac{1}{1 - \beta_1} \approx 50 \text{ steps (slow: tracks true gradient)} \quad (8)$$

$$\tau_{\text{error}} = \frac{1}{1 - \beta_2} \approx 12 \text{ steps (fast: reacts to errors)} \quad (9)$$

Intuition: Error correction must adapt *faster* than the error itself. If $\beta_2 \approx \beta_1$, both move at same timescale $\implies$ always chasing, never catching.

This is hierarchical adaptive control: fast inner loop stabilizes, slow outer loop tracks.

3.4 Comparison with Adan

Table 1: IRIS vs Adan: Algorithmic Comparison

Property	Adan	IRIS
Core quantity	$d_t = \mathbf{g}_t - \mathbf{g}_{t-1}$	$\mathbf{I}_{t t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1}$
Interpretation	Gradient velocity	Prediction error
Variance	$2\sigma^2$ (two samples)	σ^2 (vs smooth estimate)
Near-minima behavior	$d_t \rightarrow 0$ (correction vanishes)	$\mathbf{I}_{t t-1} < 0$ (active damping)
Stored gradients	Requires $\mathbf{g}_{t-1}$	No gradient storage
Memory buffers	4 ($\mathbf{m}, \mathbf{v}, \mathbf{d}, \mathbf{g}_{t-1}$)	3 ($\widehat{\mathbf{g}}, \widehat{\mathbf{v}}, \widehat{\epsilon}$)
Bias correction	Standard (β^t)	Kalman gain (recursive)
Stability	Convergence-divergence cycles	Unconditional stability

Key difference: Adan’s gradient difference becomes *less informative* as training progresses (approaches zero near minima). IRIS’s innovation remains *informative throughout*—near minima, lagging momentum creates large negative innovation that automatically damps updates.

4 Theoretical Properties

4.1 Variance Reduction and Stability

Theorem 1 (Innovation Variance Reduction). *Under stochastic gradient noise with $\mathbb{E}[\xi_t] = 0$, $\text{Var}[\xi_t] = \sigma^2$, and weak temporal correlation $|\text{Cov}[\xi_t, \xi_{t-1}]| \leq \rho\sigma^2$ where $\rho \ll 1$:*

$$\text{Var}[\mathbf{I}_{t|t-1}] \leq (1 + \epsilon(\beta_1))\sigma^2 \ll 2\sigma^2 \approx \text{Var}[d_t] \quad (10)$$

where $\epsilon(\beta_1) \rightarrow 0$ as $\beta_1 \rightarrow 1$.

Proof sketch. Expand $\widehat{\mathbf{g}}_{t-1} = \sum_{i=0}^{t-2} \beta_1^i (1 - \beta_1) \mathbf{g}_{t-1-i}$. Since $\mathbf{g}_i$ are weakly correlated:

$$\text{Var}[\widehat{\mathbf{g}}_{t-1}] = \sum_{i=0}^{t-2} [\beta_1^i (1 - \beta_1)]^2 \text{Var}[\mathbf{g}_{t-1-i}] + \text{cross terms} \quad (11)$$

$$\leq (1 - \beta_1)^2 \sigma^2 \sum_{i=0}^{\infty} \beta_1^{2i} + O(\rho) \quad (12)$$

$$= \frac{(1 - \beta_1)^2}{1 - \beta_1^2} \sigma^2 + O(\rho) = \frac{1 - \beta_1}{1 + \beta_1} \sigma^2 + O(\rho) \quad (13)$$

For $\beta_1 = 0.98$: $\text{Var}[\widehat{\mathbf{g}}_{t-1}] \approx 0.01\sigma^2$. Thus $\text{Var}[\mathbf{I}_{t|t-1}] \approx \sigma^2 + 0.01\sigma^2 \approx \sigma^2$. $\square$

Corollary 1 (Stability at High Variance Smoothing). *With $\beta_3 = 0.999$ (high variance smoothing), IRIS maintains stable convergence while Adan oscillates due to the $2\times$ variance in d_t overwhelming the smoothing.*

4.2 Near-Minima Automatic Damping

Proposition 2 (Self-Stabilization Near Minima). *Near a local minimum where $\|\mathbf{g}_t\| \rightarrow 0$ but $\|\widehat{\mathbf{g}}_{t-1}\| > 0$ (lagging momentum):*

$$\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1} < 0 \implies \beta_2 \widehat{\epsilon}_t < 0 \quad (14)$$

Thus the update numerator $\widehat{\mathbf{g}}_t + \beta_2 \widehat{\epsilon}_t$ is automatically damped without manual learning rate decay.

Contrast with Adan: Near minima, $\mathbf{g}_t \approx \mathbf{g}_{t-1} \approx 0 \implies d_t \approx 0$. Correction vanishes exactly when stabilization is needed most, requiring manual learning rate schedules.

4.3 Learning Rate Tolerance

Proposition 3 (Extreme LR Robustness). *Under Lipschitz gradients ($\|H\| \leq L$), IRIS with $\beta_3 = 0.999$ maintains bounded updates for learning rates up to $\eta \sim O(1/L)$, while Adan and Adam require $\eta \sim O(1/(10L))$ and $O(1/(100L))$ respectively.*

Mechanism: Three synergistic stabilizers:

1. **Innovation damping:** Large steps $\implies$ large $|\mathbf{I}_{t|t-1}| \implies$ negative correction
2. **Variance matching:** $\widehat{\mathbf{v}}_t \propto (\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2$ increases with step size $\implies$ self-regulation
3. **High β_3 :** Slow variance adaptation prevents denominator collapse

4.4 Convergence Analysis

Theorem 2 (Convergence Rate). *Under standard assumptions (Lipschitz gradients, bounded variance, smooth f), IRIS achieves:*

- **Convex:** $O(1/\sqrt{T})$ convergence in expectation
- **Non-convex:** $\min_{t \leq T} \mathbb{E}[\|\nabla f(\mathbf{x}_t)\|^2] = O(1/\sqrt{T})$

Proof sketch. Standard analysis for adaptive gradient methods applies. Key: bias-corrected moments ensure $\mathbb{E}[\widehat{\mathbf{g}}_t] = \mathbb{E}[\mathbf{g}_t]$ and variance estimate $\widehat{\mathbf{v}}_t$ concentrates around $\mathbb{E}[(\mathbf{g}_t + \beta_2 \mathbf{I}_{t|t-1})^2]$. Recursive bias correction (ψ updates) provides exact correction from step 1, eliminating warmup artifacts. Full proof in Appendix A.2. $\square$

5 Experimental Evaluation

5.1 Experimental Setup

Optimizers compared:

- **IRIS:** $\beta_1 = 0.98, \beta_2 = 0.92, \beta_3 = 0.999, \lambda = 0.01, \varepsilon = 10^{-8}$
- **Adan:** Official implementation [Xie et al., 2024], tuned $(\beta_1, \beta_2, \beta_3)$
- **Adam, RAdam, NAdam:** PyTorch defaults
- **AdamW:** PyTorch default with decoupled weight decay

Hardware: [Your GPU specs], PyTorch 2.0+, CUDA 11.8

5.2 Toy Function Benchmarks

We evaluate on classic optimization test functions to isolate algorithmic behavior without confounding factors (architecture, data, etc.).

5.2.1 Rosenbrock: Narrow Valley Navigation

Rosenbrock function $f(x, y) = (1 - x)^2 + 100(y - x^2)^2$ has a narrow parabolic valley ($\kappa \sim 100$) that challenges optimizers.

Table 2: Learning Rate Robustness on Rosenbrock (1000 steps)

Optimizer	LR	Final Distance	Converged?	Spikes?
IRIS	1.0	0.07	✓	No
IRIS	0.3	0.05	✓	No
IRIS	0.1	0.05	✓	No
Adan	0.2	0.10	✓	Yes (step 400)
Adan	0.1	0.40	✓	Yes (step 600)
Adan	1.0	>100	×	Catastrophic
Adam	0.01	200	×	-
Adam	0.001	200	×	-
RAdam	0.001	0.6	Partial	-
RAdam	0.01	>100	×	-
NAdam	0.01	200	×	-

Key findings:

- **IRIS:** Stable across LR $\in [0.1, 1.0]$ (10× range), best distance 0.05
- **Adan:** Requires LR ≤ 0.15 (7× narrower), exhibits spikes even at LR=0.2
- **Adam:** Requires LR ≤ 0.01 (100× narrower), fails to converge
- **RAdam/NAdam:** Similar to Adam, completely fail at practical learning rates

5.2.2 Rastrigin: Local Minima Escape

Rastrigin function $f(\mathbf{x}) = 10d + \sum_i [x_i^2 - 10 \cos(2\pi x_i)]$ has many local minima.

Key finding: IRIS is the **only optimizer to reach the global minimum** (distance 10^{-6}). All competitors either:

- Trapped in local minima (Adam/RAadam/NAadam at distance ~ 1)
- Chaotic with catastrophic spikes (Adan: $10^{-4} \rightarrow 10^2$ at steps 600, 800)

Mechanism: Innovation $\mathbf{I}_{t|t-1} = \mathbf{g}_t - \hat{\mathbf{g}}_{t-1}$ provides exploratory signal even in flat regions. When stuck in local minimum ($\mathbf{g}_t \approx 0$), lagging momentum $\hat{\mathbf{g}}_{t-1} > 0$ creates negative innovation that helps escape.

5.3 CIFAR-100: Real-World Performance

Setup: ResNet-18 on CIFAR-100, batch size 2048, 200 epochs, SGD-style augmentations.

Table 3: CIFAR-100 Results (ResNet-18, Batch 2048, 200 Epochs)

Optimizer	LR	Warmup	Final Acc	Train Time
AdamW	3e-3	20 epochs	73.0%	150 min
AdaBelief	3e-3	20 epochs	71.5%	~150 min
Adan	0.003-0.01	10-20 epochs	~73%	Failed
IRIS	3e-3	None	73.8%	113 min

Key results:

- **+0.8% accuracy** over AdamW (73.8% vs 73.0%)
- **25% faster training** (113 vs 150 minutes)
- **No warmup required** (vs 20 epochs for AdamW/AdaBelief)
- **AdaBelief overfits:** Lower training loss but -2% test accuracy (71.5% vs 73.8%)
- **Adan diverged:** Extensive tuning (6+ configs) failed to match baselines

Training ResNet-18 on CIFAR-100




Figure 1: **ResNet-18 on CIFAR-100.** Test loss (left) and test accuracy (right) comparing IRIS (labeled Nova in plot) against AdamW. IRIS reaches lower test loss and higher final accuracy without warmup.




Figure 2: **Test loss comparison** (EMA-smoothed) for IRIS vs AdamW over 200 epochs.




Figure 3: **Second-moment statistics** (max and median) over training. IRIS (Nova) maintains a healthier second-moment profile than AdamW, consistent with variance matching of the corrected gradient.

5.4 Ablation Studies

Table 4: Component Ablations (Rosenbrock, LR=1.0)

Configuration	Final Distance	Stable?
IRIS (full)	0.07	✓
IRIS ($\beta_2 = 0$, no error term)	0.15	✓
IRIS ($\beta_2 = \beta_1$, no timescale separation)	0.20	Oscillatory
IRIS ($\beta_3 = 0.99$, lower smoothing)	0.10	Oscillatory
IRIS (Adan’s gradient difference)	>100	×

Key findings:

- Removing error term ($\beta_2 = 0$): Still works but 2× worse (0.15 vs 0.07)
- No timescale separation ($\beta_2 = \beta_1$): Oscillatory, worse convergence
- Lower variance smoothing ($\beta_3 = 0.99$): Oscillations appear (but still converges!)
- Using gradient difference: **Catastrophic failure** (>100 distance)

5.5 Computational Efficiency

Table 5: Computational Cost Analysis

Optimizer	Mem Buffers	Vector Ops/Step	Time/Step
Adam	2	5	1.00×
AdamW	2	5	1.00×
Adan	4	7	1.05×
IRIS	3	6	1.02×

IRIS is essentially free: 2% overhead vs Adam, 25% less memory than Adan, while providing 10-1000× LR robustness and unconditional stability.

6 Discussion

6.1 Why Innovation Succeeds Where Gradient Difference Fails

The fundamental issue with gradient difference methods like Adan is **comparing two noisy samples**:

$$d_t = \mathbf{g}_t - \mathbf{g}_{t-1} = (\nabla f + \xi_t) - (\nabla f + \xi_{t-1}) = \underbrace{\nabla f_t - \nabla f_{t-1}}_{\text{signal}} + \underbrace{\xi_t - \xi_{t-1}}_{\text{noise}} \quad (15)$$

The noise term has variance $\text{Var}[\xi_t - \xi_{t-1}] \approx 2\sigma^2$, overwhelming the signal. No amount of variance smoothing (β_3) can fix this—the underlying correction signal is fundamentally noisy.

IRIS’s innovation compares to *smooth estimate*:

$$\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1} = (\nabla f + \xi_t) - \nabla f_{\text{smooth}} = \text{signal} + \underbrace{\xi_t}_{\sigma^2} \quad (16)$$

The noise is halved, enabling high variance smoothing ($\beta_3 = 0.999$) to produce hyperstable convergence.

6.2 The Convergence-Divergence Problem

Adan’s most critical failure mode is **convergence followed by catastrophic divergence**:

- **Rosenbrock:** Converges to 10^{-4} by step 200, then spikes to 10^{-1} at step 400 (10,000× increase)
- **Rastrigin:** Repeated spikes of $10^6\times$ at steps 600, 800

Why this happens: Near minima, $\mathbf{g}_t \approx \mathbf{g}_{t-1} \approx 0 \implies d_t \approx 0$. Correction vanishes while momentum remains large. Small perturbations cause unchecked momentum explosion.

Why IRIS avoids this: Innovation $\mathbf{I}_{t|t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1}$ remains informative. Near minima, $\mathbf{g}_t \rightarrow 0$ but $\widehat{\mathbf{g}}_{t-1} > 0$ (lagging) creates *large negative innovation* that actively damps updates.

This makes Adan unusable for:

- **Checkpointing:** May save model during spike
- **Early stopping:** May stop during recovery phase
- **LR scheduling:** Spikes confuse schedulers
- **Distributed training:** Gradient explosions cause synchronization failures

IRIS’s unconditional stability eliminates all these failure modes.

6.3 Practical Implications

Hyperparameter tuning elimination: IRIS’s extreme LR robustness (works across $\text{LR} \in [0.3, 1.0]$) means practitioners can:

1. Use $\text{LR} = 1.0$ as universal default

2. Skip grid search entirely
3. Transfer learning rates across tasks

This reduces experimentation cost from $O(10)$ runs (grid search) to $O(1)$ (single run).

No warmup needed: Recursive bias correction provides exact correction from step 1, eliminating the 10-20 epoch warmup required by Adam/AdamW/Adan.

Training acceleration: CIFAR-100 speedup (113 vs 150 min) comes from:

- Innovation reuse (compute once, use 3 times)
- Fused kernels (`lerp_()`)
- No GPU-CPU sync (recursive bias correction)

6.4 Limitations

- **Additional hyperparameter:** β_2 must be tuned relative to β_1 (recommended: $\beta_2 = \beta_1 - 0.2$). However, this is *far less sensitive* than learning rate.
- **Not always best:** On carefully-tuned Adam+schedule baselines, IRIS may match but not exceed performance. IRIS’s advantage is *eliminating the need for careful tuning*.
- **Theory-practice gap:** While we provide convergence guarantees, the extreme LR tolerance (1000 $\times$ vs Adam) exceeds what theory currently predicts. Tighter bounds remain open.

6.5 Future Work

- **Adaptive β scheduling:** Can β_1 increase during training (Kalman-inspired)?
- **Layer-wise adaptation:** Different β values per layer?
- **Second-order extensions:** Combine innovation with Hessian information?
- **Theoretical tightening:** Prove 1000 $\times$ LR advantage formally

7 Conclusion

We introduced IRIS, an innovation-based optimizer that achieves 10-1000 $\times$ higher learning rate tolerance than existing methods while maintaining unconditional stability. By tracking prediction error ($\mathbf{l}_t|_{t-1} = \mathbf{g}_t - \widehat{\mathbf{g}}_{t-1}$) rather than gradient velocity ($d_t = \mathbf{g}_t - \mathbf{g}_{t-1}$), IRIS reduces correction variance by 50% and enables automatic near-minima damping.

Our key contributions:

1. **Theoretical:** Proved innovation has 50% lower variance than gradient difference (Theorem 1)
2. **Algorithmic:** Introduced Kalman-inspired recursive bias correction, eliminating warmup
3. **Empirical:** Demonstrated extreme LR robustness (Rosenbrock LR=1.0), local minima escape (Rastrigin), and real-world gains (CIFAR-100: +0.8% accuracy, 25% faster)

IRIS represents a paradigm shift: from *hyperparameter-sensitive optimization requiring extensive tuning* to *universally-robust optimization working out-of-the-box*. A single learning rate (0.3-1.0) works across diverse tasks, eliminating the grid search that plagues deep learning practice.

Code: <https://github.com/madhavasthana-oss/iris>

Acknowledgments

[Your acknowledgments here]

References

- John Duchi, Elad Hazan, and Yoram Singer. Adaptive subgradient methods for online learning and stochastic optimization. In *Journal of Machine Learning Research*, volume 12, pages 2121–2159, 2011.
- Vineet Gupta, Tomer Koren, and Yoram Singer. Shampoo: Preconditioned stochastic tensor optimization. In *International Conference on Machine Learning (ICML)*, 2018.
- Rudolph Emil Kalman. A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1):35–45, 1960.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *International Conference on Learning Representations (ICLR)*, 2015.
- Dong C. Liu and Jorge Nocedal. On the limited memory BFGS method for large scale optimization. *Mathematical Programming*, 45(1):503–528, 1989.
- Hong Liu, Zhiyuan Li, David Hall, Percy Liang, and Tengyu Ma. Sophia: A scalable stochastic second-order optimizer for language model pre-training. In *International Conference on Learning Representations (ICLR)*, 2024.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations (ICLR)*, 2019.
- James Martens and Roger Grosse. Optimizing neural networks with Kronecker-factored approximate curvature. In *International Conference on Machine Learning (ICML)*, 2015.
- Yurii Nesterov. A method for solving the convex programming problem with convergence rate $O(1/k^2)$. *Soviet Mathematics Doklady*, 27:372–376, 1983.
- Laurent Olivier et al. Kalman-based stochastic gradient method with application to prognosis and optimization of batch processes. *arXiv preprint*, 2021.
- Sashank J. Reddi, Satyen Kale, and Sanjiv Kumar. On the convergence of Adam and beyond. In *International Conference on Learning Representations (ICLR)*, 2018.
- Tijmen Tieleman and Geoffrey Hinton. Lecture 6.5—RmsProp: Divide the gradient by a running average of its recent magnitude. COURSE: Neural Networks for Machine Learning, 2012.
- James Vuckovic. Kalman gradient descent: Adaptive variance reduction in stochastic optimization. *arXiv preprint arXiv:1810.12273*, 2018.
- Xingyu Xie, Pan Zhou, Huan Li, Zhouchen Lin, and Shuicheng Yan. Adan: Adaptive nesterov momentum algorithm for faster optimizing deep models. In *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 2024.
- Zhewei Yao, Amir Gholami, Sheng Shen, Mustafa Mustafa, Kurt Keutzer, and Michael Mahoney. AdaHessian: An adaptive second order optimizer for machine learning. In *AAAI Conference on Artificial Intelligence*, 2021.

Juntang Zhuang, Tommy Tang, Yifan Ding, Sekhar Tatikonda, Nicha Dvornek, Xenophon Papademetris, and James Duncan. Adabelief optimizer: Adapting stepsizes by the belief in observed gradients. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

A Detailed Proofs

A.1 Proof of Theorem 1

A.2 Proof of Theorem 2

We begin by establishing coefficients and lemmas on coefficients the first constraint we impose is

$$\beta_1 \geq \beta_2$$

And the following coefficients

$$\alpha_t \leftarrow \frac{1 - \beta_1}{1 - \beta_1^t} + \beta_2 \frac{1 - \beta_2}{1 - \beta_2^t} \quad (17)$$

$$\gamma_t \leftarrow \beta_1 \frac{1 - \beta_1^{t-1}}{1 - \beta_1^t} - \beta_2 \frac{1 - \beta_2}{1 - \beta_2^t} \quad (18)$$

$$\delta_t \leftarrow \beta_2^2 \frac{1 - \beta_2^{t-1}}{1 - \beta_2^t} \quad (19)$$

Where, if $\mathbf{G}_{t,i} = \widehat{\mathbf{g}}_{t,i} + \beta_2 \widehat{\mathbf{e}}_{t,i}$ Notice also that $\alpha_t + \gamma_t = 1$, and from the formulation of these coefficients, $\mathbf{G}_{t,i} \leftarrow \alpha_t \mathbf{g}_{t,i} + \gamma_t \widehat{\mathbf{g}}_{t-1,i} + \delta_t \widehat{\mathbf{e}}_{t-1,i}$ Then, if regret gets bounded as

$$f(\mathbf{x}_t) - f(\mathbf{x}^*) \leq \mathbf{g}_t^\top (\mathbf{x}_t - \mathbf{x}^*) \leq \sum_i \mathbf{g}_{t,i} (\mathbf{x}_{t,i} - \mathbf{x}_i^*)$$

then the term being summed gets bounded as

$$\mathbf{g}_{t,i} (\mathbf{x}_{t,i} - \mathbf{x}_i^*) \leq \frac{(\mathbf{x}_{t,i} - \mathbf{x}_i^*)^2 - (\mathbf{x}_{t+1,i} - \mathbf{x}_i^*)^2}{2\eta_t \alpha_t} \sqrt{\widehat{\mathbf{v}}_{t,i}} + \frac{\eta_t}{2\alpha_t} \frac{\mathbf{G}_{i,t}^2}{\sqrt{\widehat{\mathbf{v}}_{i,t}}} - \frac{\mathbf{x}_{t,i} - \mathbf{x}_i^*}{2\eta_t \alpha_t} (\gamma_t \widehat{\mathbf{g}}_{t-1,i} + \delta_t \widehat{\mathbf{e}}_{t-1,i})$$

To proceed, we show that $\|\widehat{\mathbf{e}}_t\|_\infty \leq \Lambda_\infty$ and $\|\widehat{\mathbf{e}}_t\|_2 \leq \Lambda$. A simple proof is as follows,

$$\begin{aligned} \widehat{\mathbf{e}}_t &= \frac{\sum_{i=1}^t \beta_2^{t-i} (\mathbf{g}_i - \widehat{\mathbf{g}}_{i-1})}{\sum_{i=1}^t \beta_2^{t-i}} \\ &= \frac{\sum_{i=1}^t \beta_2^{t-i} \left(\mathbf{g}_i - \frac{\sum_{j=1}^{i-1} \beta_1^{i-j-1} \mathbf{g}_j}{\sum_{j=1}^{i-1} \beta_1^{i-j-1}} \right)}{\sum_{i=1}^t \beta_2^{t-i}} \end{aligned}$$

B Additional Experimental Results

C Hyperparameter Sensitivity Analysis

D Implementation Details